

ASST3: Virtual Memory

Table of Contents

- [Checklist](#)
- [Due Dates and Mark Distribution](#)
- [Introduction](#)
 - [The System/161 TLB](#)
 - [The System/161 Virtual Address Space Map](#)
- [Setting Up Assignment 3](#)
 - [Configure OS/161 for Assignment 3](#)
 - [Building for ASST3](#)
- [Coding Assignment](#)
 - [Address Space Management](#)
 - [Address Translation](#)
 - [Testing and Debugging Your Assignment](#)
 - [Hints](#)
- [FAQ, Gotchas and Video](#)
- [Basic Assignment Submission](#)
- [Advanced Assignment](#)
 - [Advanced Assignment Submission](#)

Checklist

- Read this spec.
- Set up the assignment.
- Watch the assignment overview video.
- Do the week 9 tutorial exercises yourself.
- Attend the tutorial.
- Read the FAQ for assignment 3 on the wiki.
- Do the assignment.
- Don't forget to commit your most recent changes.
- Submit your assignment.

Due Dates and Mark Distribution

Due Date & Time: 16:00 Friday 21st April

Marks: The base assignment is worth 30 marks (of the 100 available for the class mark component of the course)

The 2% per day bonus for each day early applies, capped at 10%, as per course outline.

Students can submit the advanced part with approval. Students obtain approval if they submit a solution to the basic assignment 5 days prior to the due date. Marks obtained are added to any shortfall in the class mark component up to a maximum of 5 bonus marks for this assignments, and 10 marks overall for all across all assignments.

There are familiarisation questions in your week 9 tutorial. Please answer the questions and bring them to your tutorial. Feel free to ask any assignment related questions need to clarify your understanding.

Introduction

In this assignment you will implement the virtual memory sub-system of OS/161. The existing VM implementation in OS/161, *dumbvm*, is a minimal implementation with a number of shortcomings. In this assignment you will adapt OS/161 to take full advantage of the simulated hardware by implementing address space management, page table management, and management of the MIPS software-managed Translation Lookaside Buffer (TLB).

The System/161 TLB

This section provides a summary of the R3000 (System/161) virtual memory mechanisms. Further info is provided in the lectures and in the *R3000 Reference Manual* and *Hardware Guide* on the [course website](#).

The R3000 TLB entry includes a 20-bit virtual page number and a 20-bit physical frame number as well as the following five fields:

- global: 1 bit; if set, ignore the PID bits in the TLB.
- valid: 1 bit; set if the TLB entry contains a valid translation.
- dirty: 1 bit; enables writing to the page referenced by the entry; if this bit is 0, the page is only accessible for reading.
- nocache: 1 bit; unused in System/161. In a real processor, indicates that the hardware cache will be disabled when accessing this page.
- asid: 6 bits; a context or address space ID that can be used to allow entries to remain in the TLB after a context switch.

All these bits/values are maintained by the operating system (i.e. your code). When the valid bit is set, the TLB entry contains a valid translation. This implies that the virtual page is present in physical memory. A TLB miss occurs when no TLB entry can be found with a matching virtual page and address space ID (unless the global bit is set in which case the address space ID is ignored) and a valid bit that is set.

For this assignment, you may largely ignore the ASID field and set it to zero in your TLB entries. Note: In OS/161, `as_activate()` is called whenever a new address space becomes active in the TLB, so `as_activate()` is typically programmed to flush the TLB (why?).

The System/161 Virtual Address Space Map

The MIPS divides its address space into several regions that have hardwired properties. These are:

- kseg2, TLB-mapped cacheable kernel space
- kseg1, direct-mapped uncached kernel space
- kseg0, direct-mapped cached kernel space
- kuseg, TLB-mapped cacheable user space

Both direct-mapped segments map to the first 512 megabytes of the physical address space.

The top of kuseg is 0x80000000. The top of kseg0 is 0xa0000000, and the top of kseg1 is 0xc0000000.

The memory map thus looks like this:

Address	Segment	Special Properties
0xffffffff	kseg2	
0xc0000000		
0xbfffffff	kseg1	
0xbfc00180		Exception address if BEV set.
0xbfc00100		UTLB exception address if BEV set.
0xbfc00000		Execution begins here after processor reset.
0xa0000000		

0x9fffffff	kseg0	
0x80000080		Exception address if BEV not set.
0x80000000		UTLB exception address if BEV not set.
0x80000000		
0x7fffffff	kuseg	
0x00000000		

Setting Up Assignment 3

We assume after ASST0, ASST1, and ASST2 that you now have some familiarity with setting up for OS/161 development. If you need more detail, refer back to ASST0.

Clone the ASST3 source repository from gitlab.cse.unsw.edu.au. Note: replace XXX with your 3 digit group number.

```
% cd ~/cs3231
% git clone https://zNNNNNNN@gitlab.cse.unsw.edu.au/COMP3231/23T1/grpXXX-asst3.git asst3-src
```

Note: The gitlab repository is shared between you and your partner. You can both push and pull changes to and from the repository to cooperate on the assignment.

Configure OS/161 for Assignment 3

Remember to set your `PATH` environment variable as in previous assignments (or run the `3231` command).

Before proceeding further, configure your new sources, and build and install the user-level libraries and binaries.

```
% cd ~/cs3231/asst3-src
% ./configure
% bmake
% bmake install
```

You have to reconfigure your kernel before you can use the framework provided to do this assignment. The procedure for configuring a kernel is the same as before, except you will use the ASST3 configuration file:

```
% cd ~/cs3231/asst3-src/kern/conf
% ./config ASST3
```

You should now see an `ASST3` directory in the `compile` directory.

Building for ASST3

When you built OS/161 for ASST0, you ran `bmake` from `compile/ASST0`. When you built for ASST1, you ran `bmake` from `compile/ASST1` ... you can probably see where this is heading:

```
% cd ../compile/ASST3
% bmake depend
% bmake
% bmake install
```

If you now run the kernel as you did for previous assignments, you should get to the menu prompt. If you try and run a program, it will fail with a message about an unimplemented feature (the failure is due to the unimplemented `as_*` functions that you must write). For example, run `p /bin/true` at the OS/161 prompt to run the program `/bin/true` in `~/cs3231/root`.

```
OS/161 kernel [? for menu]: p /bin/true
Running program /bin/true failed: Function not implemented
Program (pid 2) exited with status 1
Operation took 0.173469806 seconds
OS/161 kernel [? for menu]:
```

Note: If you don't have a `sys161.conf` file, you can use the one from ASST1.

The simplest way to install it is as follows:

```
% cd ~/cs3231/root
% wget http://cgi.cse.unsw.edu.au/~cs3231/22T1/assignments/asst1/sys161.conf -O sys161.co
nf
```

You are now ready to start the assignment.

Coding Assignment

This assignment involves designing and implementing a number of data-structures and the functions that manipulate them. Before you start, you should work out what data you need to keep track of, and what operations are required.

Address Space Management

OS/161 has an address space data type that encapsulates the book-keeping needed to describe an address space: the `struct addrspace`. To enable OS/161 to interact with your VM implementation, you will need to implement the functions in `kern/vm/addrspace.c` and potentially modify the data type. The semantics of these functions is documented in `kern/include/addrspace.h`.

Note: You may use a fixed-size stack region (say 16 pages) for each process.

Address Translation

The main goal for this assignment is to provide virtual memory translation for user programs. To do this, you will need to implement a page table data structure and its associated TLB refill handler. For this assignment, you will implement a **hashed page table (HPT)**.

Note that a hashed page table is a fixed sized data structure allocated at boot time and shared between all processes. We suggest sizing the table to have twice as many entries as there are frames of physical memory in RAM. The size of physical memory can be obtained via a call to `ram_getsize()`. You can allocate and initialise your HPT together with any other associated variables in `vm_bootstrap()` in `vm.c`.

Given the HPT is a shared data structure, it will have to handle concurrent access by multiple processes. You'll need to synchronise operations that access to avoid potential races. Using single lock covering the entire HPT to ensure all operations execute mutually exclusively is satisfactory approach.

Each entry in the HPT typically has a process identifier, the page number, a link to handle collisions, and a frame number and permissions in `EntryLo` format for faster TLB loading. A HPT entry should not need to exceed 4 32-bit words.

One can use the OS/161 address space pointer (of type `struct addrspace`) as the value of the current process ID. It is readily accessible where needed, and is unique to each address space.

A simple hash function that produces an index into the HPT can be derived from the following example code.

```
index = (((uint32_t) addrspace_ptr) ^ (faultaddr >> 12)) % hpt_size;
```

Note that `(faultaddr >> 12)` is the page number of the entry we are looking for.

Hash collisions are always possible even with a good hash function, and you can use either internal or external chaining to resolve collisions. We suggest external chaining to avoid the HPT filling in the presence of sharing (in the advanced/bonus assignments), however internal chaining is sufficient for the basic assignment.

Note: For the basic assignment, a HPT using internal chaining should not run out of available slots in our testing, assuming a HPT with at least twice as many entries as frames in RAM. It is sufficient for the basic assignment to return `ENOMEM` if your internally-chained HPT runs out of slots.

The following questions may assist you in designing the contents of your page table

- What information do you need to store for each page?
- How does the page table get populated?
- When are frames allocated to back pages.

Note: Applications expect pages to contain zeros when first used. **This implies that newly allocated frames that are used to back pages should be zero-filled prior to mapping**

Testing and Debugging Your Assignment

To test this assignment, you should run a process that requires more virtual memory than the TLB can map at any one time. You should also ensure that touching memory not in a valid region will raise an exception. The huge and faulter tests in `testbin` may be useful. See the Wiki for more options.

Apart from GDB, you may also find the `trace161` command useful. `trace161` will run the simulator with tracing, for example

```
% trace161 -t t -f outfile kernel
```

will record all TLB accesses in `outfile`.

Don't use `kprintf()` for `vm_fault()` debugging. See the Wiki for more info.

Hints

One approach to implementing the assignment is in the following order:

- Review how the specified page table works from the lectures, and understand its relationship with the TLB.
- Review the assignment specification and its relationship with the supplied code.
 - *dumbvm* is not longer compiled into the OS/161 kernel for this assignment (`kern/arch/mips/vm/dumbvm.c`), but you can review it as an example implementation within the interface/framework you will be working within.
 - The only candidate for code re-use is the TLB flush in `as_activate()`.
 - Note: Your implementation of TLB refill in `vm_fault()` should use `tlb_random()`.
- Work out a basic design for your page table implementation.
- Modify `kern/vm/vm.c` to insert , lookup, and update page table entries in your page table structure.
- Implement the TLB exception handler `vm_fault()` in `vm.c` to refill the TLB and keep it consistent with your page table.
- Implement the functions in `kern/vm/addrspace.c` that are required for basic functionality (e.g. `as_create()`, `as_prepare_load()`, etc.). Allocating user pages in `as_define_region()` may also simplify your assignment, however good solution allocate pages in `vm_fault()`.
 - E.g. `as_create()` should initialise your page table, `as_destroy()` should clean it up.
- Test and debug this. Use the debugger or `trace161`!

Note: Interrupts should be disabled when writing to the TLB, see `dumbvm` for an example. Otherwise, unexpected concurrency issues can occur.

`as_activate()` and `as_deactivate()` can be copied from `dumbvm`.

FAQ, Gotchas and Video

Don't forget to look at <https://wiki.cse.unsw.edu.au/cs3231cgi/2023t1/Asst3> for an up to date list of potential issues you might encounter.

There is also an overview video on the assignment available on the lectures page in the course account <https://cgi.cse.unsw.edu.au/~cs3231/lectures.php>.

Basic Assignment Submission

The submission instructions are available on the [Wiki](#). Like previous assignments, you will be submitting the git repository bundle via CSE's `give` system. For ASST3, the submission system will do a test build and run a simple test to confirm your bundle at least compiles.

Warning! Don't ignore the submission system! If your submission fails the submission process, you may not receive any marks.

Warning! Don't forget to commit your changes prior to generating your bundle.

To submit your bundle:

```
% cd ~  
% give cs3231 asst3 asst3.bundle
```

You're now done.

Even though the generated bundle should represent all the changes you have made to the supplied code, occasionally students do something "ingenious". So always keep your git repository so that you may recover your assignment should something go wrong.

Advanced Assignment

The advanced assignment consists of a student-chosen subset of the problems below. The total marks available are capped at 5 marks.

Students can do the advanced part with the permission of the lecturer, and only if basic assignment is completed a week prior to the deadline.

- (easy) 2 marks Shared pages and copy-on-write.
- (easy) 2 marks Implement `sbrk()` to enable user-level `malloc()` to function with more memory than its initially allocated pool.
- (hard) 3 marks.
 - Implement a simplified `mmap()` and `munmap()`. Note: you only need to support the simplified case of mapping a file, and `munmap()` the entire region that was mapped.

The prototypes are expected to be

```
void *mmap(size_t length, int prot, int fd, off_t offset);  
int munmap(void *addr);
```

Where prot can be PROT_READ and/or PROT_WRITE. Compared to traditional mmap, there are no flags, and the OS chooses the virtual address to locate the region. You must ensure that applications can open a file, update it, and have the updated file propagate to the filesystem.

- And, implement demand-loading. You should load pages only when they are referenced by the user process, as opposed to at process creation.
- (seriously hard) 5 marks Implement paging. You should implement some page replacement algorithm and demonstrate your solution running under memory pressure.

Given you're doing the advanced version of the assignment, I'm assuming you are competent with managing your git repository and don't need detailed directions. We expect you to work on a specific branch in your repository to both build upon your existing assignment, while keeping your advanced assignment separate at the same time.

Here are some git commands that will be helpful.

- One member of your group should create the branch and push it back to gitlab.

```
git checkout -b asst3_adv
git push --set-upstream origin asst3_adv
```

- To switch back to the basic assignment at some point.

```
git checkout master
```

- To switch to the advanced assignment at another point.

```
git checkout asst3_adv
```

Advanced Assignment Submission

Submission for the advanced assignment is similar to the basic assignment, **except the advance component is given to a separate assignment name: asst3_adv**. Again, you need to generate a bundle based on your repository. Note: Our marking scripts will switch to the asst3_adv branch prior to testing the advanced assignment.

Submit your solution

```
% cd ~
% give cs3231 asst3_adv asst3_adv.bundle
```

You're now done.