

COMP3411/9814 Revision

# Exam

- 2 May 2023
- Online – Moodle
- Reading time starts 1:45pm
- Open for answers 2pm to 4pm (ELS students have extra time)
- 20 multiple choice questions (very similar to quiz questions)
- 6 multi-part written questions (max 16 sub-parts)
  - Short answer, some calculation
  - Not more complex than multiple choice, so don't panic.
- Supplementary exams in week 22-26 May (exact date not set yet)

# Search


- What are the different types of search algorithms?
- What are their properties?
  - Order of search
  - Optimality
  - Complexity
- Can you demonstrate different search algorithms on a simple example?

# Example Question

- Given Start (A) and Goal (I) nodes, can you work out the order of nodes visited in:

- A) Depth-first search
- B) Breadth-first search
- C) Iterative Deepening
- D) Uniform Cost Search
- E) A\* Search
- F) Greedy Search

Node	H
A	10
B	13
C	5
D	7
E	15
F	6
G	5
H	2
I	0




# CSP

- What are the different CSP strategies?
- e.g. forward checking, arc consistency
- Can you apply them on a small example?

# Example Question



- In map colouring, can you apply forward checking to a graph?
- If you are given a partial assignment of variables, can you work out the remaining legal values by forward checking or arc consistency?



# Neural Networks

- Given a network and an activation function, can you propagate value through a network to get an output (i.e. forward pass)?
- If you are given an incomplete network and an activation function, can you find the value for the missing weights so the network returns the expected output?
- Can you apply the perceptron training rule?

# Example question



- Can you compute the output of a multilayer perceptron, given a simple activation function?

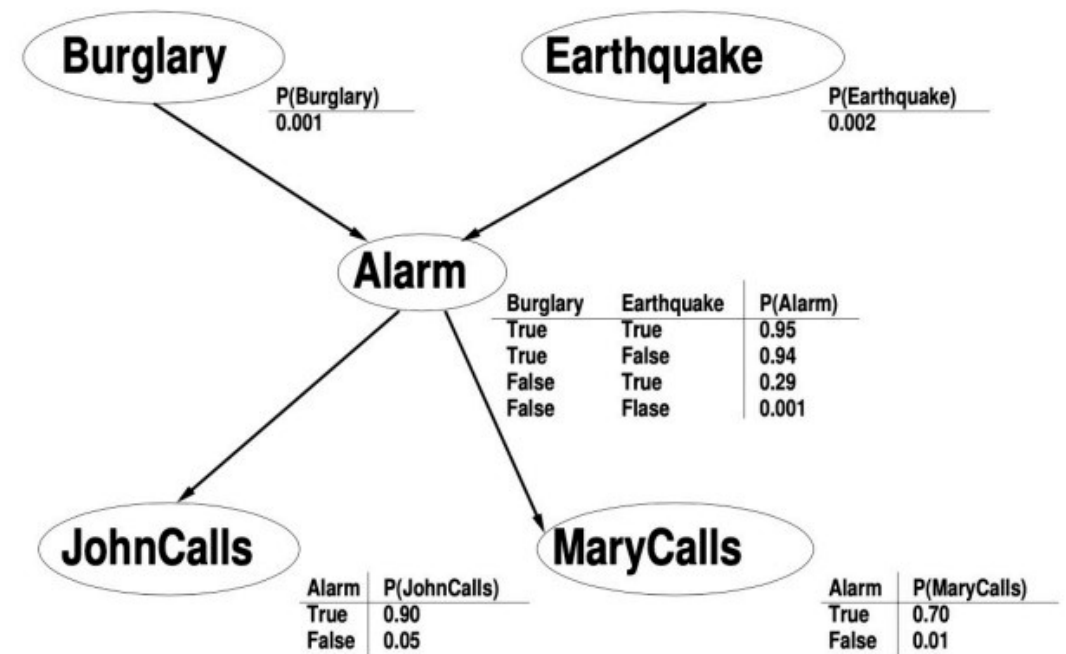


# Uncertainty

- Given a joint probability distribution table, can you find the the probability of a proposition?
- Given,  $P(A)$ ,  $P(B)$  and  $P(A|B)$ , can you calculate  $P(B|A)$ ?
- Given a Bayesian network, can you calculate different probabilities using causal inference or diagnostic inference?

# Example question

- Can you calculate  $P(E | \sim A)$ ?




# Reinforcement Learning and Robot Vision Recap

COMP3411/9814: Artificial Intelligence

# Reinforcement Learning

COMP3411/9814: Artificial Intelligence

# Reinforcement Learning



# Elements of reinforcement learning

There are four essential elements:

- **Policy**
  - Informs how to act in a particular situation.
  - Set of stimulus-response rules or associations.
  - Can be stochastic.

# Elements of reinforcement learning

There are four essential elements:

- **Reward function**
  - Defines the aim of an RL problem.
  - Maps each perceived state (or state-action pair) into a number, the reward.
  - The goal is to maximize the long-term reward.
  - In biological systems may correspond to pain and pleasure feelings.
  - Can be stochastic.

# Elements of reinforcement learning

There are four essential elements:

- **Value function**
  - Shows what is good in the long run (the reward in an immediate sense).
  - In biological systems corresponds to more refined judgments or foresight about the future from one state.
  - Actions are decided based on the value.
  - It's much harder to determine values than rewards.



# Elements of reinforcement learning

There are four essential elements:

- **Optionally, a model of the environment**
  - Imitates the environment behaviour.
  - Can predict states and reward obtained.
  - The use of models of the environment is still relatively new.

# Action-value estimation methods

## **Greedy method**

- The simplest way to choose an action: the action with the highest estimated value.
- $A_t^*$  where  $Q_t(A_t^*) = \max_a Q_t(a)$ .

## **$\epsilon$ -greedy method**

- A simple alternative: to choose the best action most of the time, and sometimes (with a small probability  $\epsilon$ ) a random one.
- $Q_t(a)$  converges to  $q_*(a)$  with probability  $1 - \epsilon$ .

# Action-value estimation methods

## Softmax method

- $\epsilon$ -greedy effectively trades off exploration and exploitation, but the selection is equitable (or fair) among actions.
- Sometimes, the worst action is very bad.
- Softmax uses action probabilities as a Boltzmann distribution.

$$\frac{e^{Q_t(a)/\tau}}{\sum_{i=1}^n e^{Q_t(i)/\tau}}$$

# Action-value estimation methods

## Softmax method

- High temperatures give almost equal probability for all actions.
- Low temperatures make a bigger difference in the probability.
- It is not clear if softmax or  $\epsilon$ -greedy performs better.
- However, softmax shows dissimilar performance making difficult the selection of the temperature.

$$\frac{e^{Q_t(a)/\tau}}{\sum_{i=1}^n e^{Q_t(i)/\tau}}$$



- Given a set of Q values, could you calculate the selection for all Q values using either  $\epsilon$ -greedy or softmax methods?







# Returns

- If the reward sequence received is  $R_{t+1}, R_{t+2}, R_{t+3}, \dots$ . We want to maximise the expected return  $G_t$ .

$$G_t = R_{t+1} + R_{t+2} + R_{t+3} + \dots + R_T$$

- In tasks with final state and that can be divided into subsequences (episodes)
- Each episode finishes in the final state and the task starts over from an initial state.
- These tasks are known as episodic tasks.

- Can you calculate a Q value based on average reward?

# Returns


- Tasks intended to be performed continuously without limit are referred to as continuous tasks (or non-episodic).
- The return could be infinite, given that  $T = \infty$ .
- In this case, the agent maximises the discounted rewards, choosing actions to maximise the discounted return:

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$


- Discount rate  $0 \leq \gamma < 1$ . Determine the present value of future rewards. If  $\gamma = 0$ , the agent is myopic. If  $\gamma \rightarrow 1$  the agent is foresighted.

- Can you calculate  $G_i$  for a give  $i$ ?


# Q-values



$r(s, a)$  (immediate reward) values




$V^*(s)$  values



$Q(s, a)$  values

$$\gamma = 0.9$$

# Grid world example



- Can you calculate the sequence a Q table of values for each action?

## Sarsa: On-Policy TD Control

- Updates after each transition from a non-terminal  $S_t$ .
- If  $S_{t+1}$  is terminal,  $Q(S_{t+1}, A_{t+1})$  is defined as zero.

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[ R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t) \right]$$

- Each element of the 5-tuple  $(S_t, A_t, R_{t+1}, S_{t+1}, A_{t+1})$  is used, this gives the name to the algorithm.
- On-policy methods continuously estimate  $q_\pi$  for policy  $\pi$ , and at the same time change  $\pi$  greedily towards  $q_\pi$ .



# Q-Learning: Off-Policy TD Control

- A simple but important breakthrough is an off-policy TD algorithm.
- The simplest way is *one-step Q-learning*:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[ R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t) \right]$$

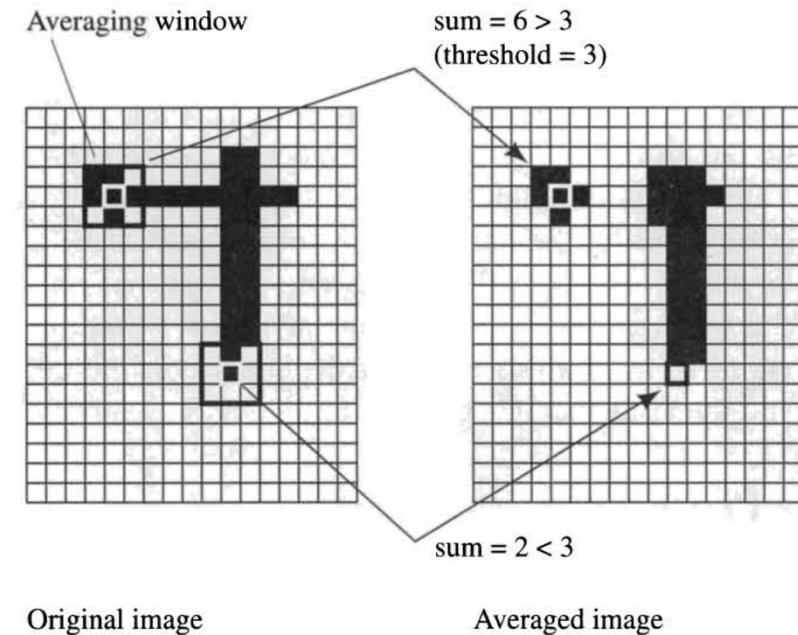
- The learned action-value function  $Q$  directly approximates  $q^*$ , the optimal action-value function, regardless the followed policy  $\pi$ .
- The policy still has an effect in which state-action pairs are visited and updated.

- Can you calculate the sequence a Q table of values for each action?
- How about for SARSA?

# Robot Vision

# Image processing: Averaging

- It can use a threshold.
- Larger rectangles achieve more smoothing.
- Broad lines are thickened and thin lines eliminated



## Image processing: Averaging

- Given a simple 4 x 4 picture matrix:


9	9	9	3
9	9	3	3
9	3	3	3
3	3	3	3

- Smooth this matrix using an averaging technique and a 3 x 3 pixel window.

## Image processing: Averaging

- There are four 3 x 3 pixel windows in the matrix.
- Replace middle value in each window by average of all the values in the window.

9	9	9	3
9	9	3	3
9	3	3	3
3	3	3	3



9	9	9	3
9	7	5	3
9	5	4	3
3	3	3	3

- Given a really small image can you demonstrate how smoothing works?

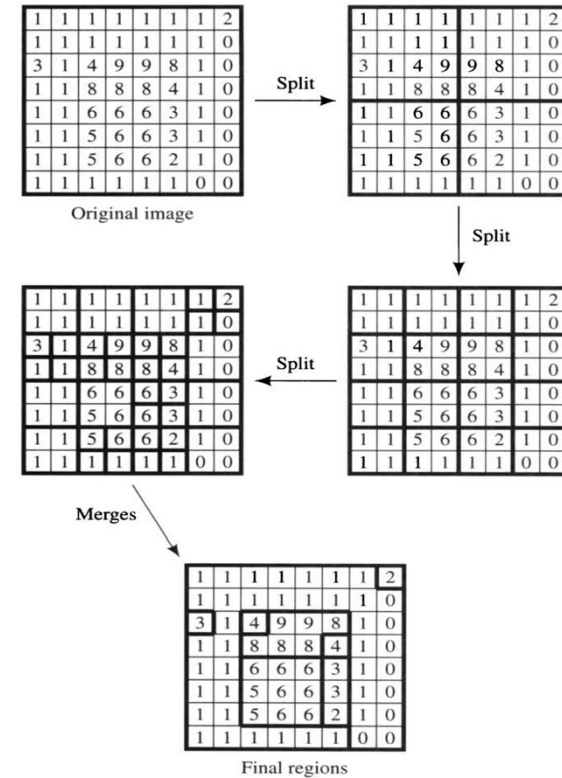


# Image processing: Region finding

- To find regions in which a property does not change abruptly.
- A region is homogeneous. Intensity difference no more than some  $\epsilon$
- Split-and-merge method.  $2^n \times 2^n$  array of pixels.
  - Each no homogeneous region is split in four.
  - Splits continues until no more splits need to be made.
  - Adjacent regions are merged if homogeneous.



# Image processing: Region finding


- Splitting and merging candidate regions.
- In this example, intensities may not vary more than 1 unit. Therefore,  $\epsilon \leq 1$ .



- Given a really small image know how region splitting and merging works?

# Stereo vision

- Under perspective projection large, distant objects might produce the same image as similar but smaller, closer ones.
- Distance estimation from single images is problematic, but sometimes possible.
- e.g., If we know an object is on the floor and the camera height.




**Figure 6.20**

Depth Calculation from a Single Image

# Stereo vision

- Depth information from stereo vision.
- Two-dimensional setup.
- Two lenses with distance  $b$ .
- Correspondence problem for pairs of points.



**Figure 6.21**

Triangulation in Stereo Vision

- Can you calculate the distance to an object from the disparity data (angles and distance between cameras)?

# Logic and Automated Reasoning

- Can you translate English sentences to and from expressions in propositional and first order logic?
- Can you work out the truth table of an expression in propositional logic?
- Do you know what soundness, completeness and decidability mean?
- Can you show something like:

$$P \wedge (Q \rightarrow R) \Leftrightarrow \neg(P \rightarrow Q) \vee (P \wedge R)$$

## Do you remember these identities?

Commutativity:

$$p \wedge q \leftrightarrow q \wedge p$$

$$p \vee q \leftrightarrow q \vee p$$

Associativity:

$$p \wedge (q \wedge r) \leftrightarrow (p \wedge q) \wedge r$$

$$p \vee (q \vee r) \leftrightarrow (p \vee q) \vee r$$

Distributivity:

$$p \wedge (q \vee r) \leftrightarrow (p \wedge q) \vee (p \wedge r)$$

$$p \vee (q \wedge r) \leftrightarrow (p \vee q) \wedge (p \vee r)$$

Implication:

$$(p \rightarrow q) \leftrightarrow (\neg p \vee q)$$

Idempotent:

$$p \wedge p \leftrightarrow p$$

$$p \vee p \leftrightarrow p$$

Double negation:

$$\neg \neg p \leftrightarrow p$$

Contradiction:

$$p \wedge \neg p \leftrightarrow \text{FALSE}$$

Excluded middle:

$$p \vee \neg p \leftrightarrow \text{TRUE}$$

De Morgan:

$$\neg (p \wedge q) \leftrightarrow (\neg p \vee \neg q)$$

$$\neg (p \vee q) \leftrightarrow (\neg p \wedge \neg q)$$

## Can you work out what the models of a formula are?

$$S \models P$$

$$S = \{p \rightarrow q, q \rightarrow p, p \vee q\}$$

$$P = p \wedge q$$

Each row is an interpretation of  $S$ .  
Only the first row is a model of  $S$ .

<b>p</b>	<b>q</b>	<b><math>p \rightarrow q</math></b>	<b><math>q \rightarrow p</math></b>	<b><math>p \vee q</math></b>	<b><math>S</math></b>	<b><math>p \wedge q</math></b>
<b>T</b>	<b>T</b>	<b>T</b>	<b>T</b>	<b>T</b>	<b>T</b>	<b>T</b>
<b>T</b>	<b>F</b>	<b>F</b>	<b>T</b>	<b>T</b>	<b>F</b>	<b>F</b>
<b>F</b>	<b>T</b>	<b>T</b>	<b>F</b>	<b>T</b>	<b>F</b>	<b>F</b>
<b>F</b>	<b>F</b>	<b>T</b>	<b>T</b>	<b>F</b>	<b>F</b>	<b>F</b>



## Proofs (propositional and first-order)

- Can you take any propositional or first order expression and turn it into conjunctive normal form?
- Given a knowledge base and a query, can you show all the steps in a resolution proof of the query?

## Prolog

- Can you write a simple Prolog program?

- E.g. merge two sorted lists:

```
?- merge([1, 3, 7, 15], [5, 9, 13], X) .  
X = [1, 3, 5, 7, 9, 13, 15]
```

## Planning

- Can you write the pre-conditions and postconditions of an action in STRIPS (PDDL) representations (i.e. preconditions and effects).
- Do you know the difference between a forward chaining and backward chaining planner?
- Given a simple planning problem, can you draw a graph showing how a partial order planner works?
- Can you explain how constraint solving can improve the efficiency of a planning algorithm?

## Inductive Logic Programming

- Given two definite clauses, can find the least general generalisation?
- Can you explain why intra\_construction, absorption and truncation are generalisation operations?

## Decision Trees

- Can you calculate the entropy resulting from split so that you can choose the most informative attribute to split on?
- Can you calculate the Laplace error and determine when to stop growing a branch in a decision tree?
- What is k-fold cross-validation and how is it used to estimate the classification accuracy of a machine learning algorithm?

## NLP

- Given a simple grammar, can you determine if an input sentence is recognised by the grammar?
- Given a simple grammar, can you generate a sentence that is recognised by the grammar?
- Can you demonstrate a left-most or right-most derivation?
- Given a definite clause grammar can you show what the output would be after recognising an input sentence?